

HMM and Part of Speech Tagging

Adam Meyers
New York University



Outline

- Parts of Speech Tagsets
- Rule-based POS Tagging
- HMM POS Tagging
- Homework Assignment



Part of Speech Tags Standards

- There is no standard set of parts of speech that is used by all researchers for all languages.
- The most commonly used English tagset is that of the Penn Treebank at the University of Pennsylvania:
 - http://repository.upenn.edu/cgi/viewcontent.cgi?article=1603&context=cis_reports
 - Page 5 of: <http://aclweb.org/anthology/J/J93/J93-2004.pdf>
 - Provides list
- To map several POS tagsets to each other, see Table 1 in:
 - <http://nlp.cs.nyu.edu/meyers/Annotation%20Compatibility%20Working%20Group%20Report%202006.html>
- POS tagsets:
 - Assume Particular Tokenizations, e.g., *Mary's* → *Mary* + *'s*
 - Distinguish inflections: e.g., *eat/VB*, *eat/VBP*, *eats/VBZ*, *ate/VBD*
 - Different instances of the same string can have different tags
 - *She wants to eat/VB*; *They eat/VBP*. *He eats/VBZ*, *Those are good eats/NNS*
- **Annotators & POS taggers assign tags to each token in a sentence, no exceptions**



The Penn Treebank II POS tagset

- Verbs: VB, VBP, VBZ, VBD, VBG, VBN
 - base, present-non-3rd, present-3rd, past, -ing, -en
- Nouns: NNP, NNPS, NN, NNS
 - proper/common, singular/plural (singular includes mass + generic)
- Adjectives: JJ, JJR, JJS (base, comparative, superlative)
- Adverbs: RB, RBR, RBS, RP (base, comparative, superlative, particle)
- Pronouns: PRP, PP\$ (personal, possessive)
- Interrogatives: WP, WP\$, WDT, WRB (compare to: PRP, PP\$, DT, RB)
- Other Closed Class: CC, CD, DT, PDT, IN, MD
- Punctuation: # \$. , : () “ ” ' ' `
- Weird Cases: FW(*deja vu*), SYM (@), LS (*1, 2, a, b*), TO (*to*), POS('s, '), UH (*no, OK, well*), EX (*it/there*)
- Newer tags: HYPH, PU



Part of Speech Tagging

- POS taggers assign 1 POS tag to each input token
 - *The/DT silly/JJ man/NN is/VBZ a/DT professor/NN ./PU*
- Different ways of breaking down POS tagging:
 - Use separate “tokenizer”, program that divides string into list of tokens – POS tagger processes output
 - Incorporate tokenizer into POS tagger
- Different ways of breaking down parsing:
 - Use separate POS tagger – output of tagger is input to parser
 - Assign POS tags as part of parsing (assumed previously)
- Accurate POS tagging is “easier” than accurate parsing
 - POS tags may be sufficient information for some tasks



Some Tokenization Rules for English

- 1) Divide at spaces and hyphens.
- 2) Divide before punctuation that is followed by: a space or the end of the line
 - Define punctuation as any non-letter/non-number:
 - ``!@#$%^&*()-_+={|}\|:;'"<, > . ? /`
 - Punctuation followed by a space, other punctuation, or at the end of line should be separated from words:
 - *...and he left.*) → *and he left . ”)*
- 3) Break off the following as separate tokens when followed by a space or end of line:
 - *'s, n't, 'd, 've, 'm, 'll, 're, ...* (a short list)
- 4) Abbreviations are exceptions to rule 2:
 - Period after abbreviations should not be separate from words
 - Most cases covered by list of 100 items (or if sentence end is known)
 - Final periods are not duplicated after abbreviations (consistency issues)
 - These periods serve 2 functions simultaneously (argument for duplication)
 - These periods occupy a single character position
 - argument against duplication – difficulty with calculating character offsets



Sentence Boundaries

- Most POS taggers assume sentence divisions
- Sample sentence splitting rules:
 - End sentence after . ? !, possibly others (:;...) Begin quotes are part of next sentence. End quotes are part of previous sentence.
 - But post-abbreviation (inc, co, ...) periods are ambiguous
 - next character is lowercase – not sentence end
 - next character is uppercase or number – possible sentence end
- Most POS taggers assume sentence boundaries are given.
- Multiple sentences within quotes are assumed separate.
 - `<S> She said, “This is the way things are. </S>`
 - `<S> This is this.</S> <S> That is that.”</S>`
- **Sentence Splitting is a potentially good Final Project Topic**
 - <http://www.aclweb.org/anthology/C12-2096>
 - <http://aclweb.org/anthology/D18-2012>
 - <https://github.com/google/sentencepiece>



Rule-based POS Tagger

- Method
 - Assign lists of potential POS tags to each word based on dictionary
 - Manual rules for Out of Vocabulary (OOV) words
 - Ex: Non-initial capital → NNP; ends in S → VBZ|NNS; default → NN|JJ; etc.
 - Apply hand-written constraints until each word has only one possible POS
- Sample Constraints:
 - 1) DT cannot immediately precede a verb
 - 2) No verb can immediately precede a tensed verb: VBZ, VBP, VBD
 - Untensed: VB (base form), VBN & VBG (past & present participles)
- Example:
 - The/DT book/{NN|VB|VBP} is/VBZ on/IN the/DT table{NN|VB|VBP}
 - The/DT book/NN is/VBZ on/IN the/DT table/NN
 - DT cannot precede VB or VBP
 - VBZ cannot be preceded by VB or VBP



Probability

- Estimate of probability of future event based on past observations

$$P(event) = \frac{\text{num of events}}{\text{num of trials}}$$

- Conditional Probability: probability of X given Y

$$P(X|Y) = \frac{P(X, Y)}{P(Y)}$$

- Examples relating to POS tags (previous examples with word N-grams):
 - Out of 200 **DT** tags, 150 of them are tagging the word **the**
 - If a word is tagged **DT**, there is a 75% chance that word is **the**
 - Example of likelihood probability
 - The POS after a **DT** is **NN** 120 times and **JJ** 60 times:
 - A word following **DT** is
 - 120/200 = 60% likely to be a singular noun (**NN**)
 - 60/200 = 30% likely to be a base adjective (**JJ**)
 - Examples of transition probability (probability of tag NN or JJ, given previous tag DT)



More Math Terminology

- N instances of a variable looked at individually:

X_1^n is the same as $\{X_1, X_2, X_3, \dots, X_n\}$ in sequence

- The product of instances of X from 1 to n

$$\prod_{i=1}^n P(X_i)$$

- Max = the maximum number in a set
- Argmax = the choice of variable values that maximizes a formula (example in 2 slides)



Probabilistic Models of POS tagging

- For tokens w_1, \dots, w_n , find the most probable corresponding sequence of possible tags t_1, \dots, t_n
 - We assume that *probable* means something like “most frequently observed in some manually tagged corpus of words”.
- Penn Treebank II (a common training corpus)
 - 1 million words from the Wall Street Journal
 - Tagged for POS (and other attributes)
- The specific sequence (sentence) is not in the training corpus
 - Therefore the actual “probability” is 0
 - Common practice: estimate probability given assumptions, e.g.,
 - Assume that we can estimate probability of whole tag sequence by multiplying simpler probabilities, e.g., sequences of 2 consecutive tags



Probabilistic Assumptions of HMM Tagging

- $\hat{t} = \underset{t_1^n}{\operatorname{argmax}} P(t_1^n | w_1^n)$
 - Choose the tag sequence of length n that is most probable given the input token sequence
- Bayes Rule:
 - $P(x|y) = \frac{P(y|x)P(x)}{P(y)}$
 - Way to derive the probability of x given y when you know: the probability of y given x, the probability of x and the probability of y
- Applying Bayes Rule to Tag Probability
 - $\hat{t} = \underset{t_1^n}{\operatorname{argmax}} \frac{P(w_1^n | t_1^n) P(t_1^n)}{P(w_1^n)}$



Simplifying Assumptions for HMMs

- Simplification: Drop the denominator
 - Denominator is same for all the tag sequences (the word sequence is given)
 - $\hat{t} = \underset{t_1^n}{\operatorname{argmax}} P(w_1^n | t_1^n) P(t_1^n)$
 - For each tag sequence calculate the product of:
 - The probability of the word sequence given the tag sequence (**likelihood**)
 - The probability of the tag sequence (**prior probability**)
 - Still too hard
- 2 simplifying assumptions make it possible to estimate the probability of tag sequences given word sequences:
 - 1) If the probability of a word is only dependent on its own POS tag,
 - $P(w_1^n | t_1^n) \approx \prod_{i=1}^n P(w_i | t_i)$
 - 2) If the probability of a POS tag is only dependent on the previous POS tag,
 - $P(t_1^n) \approx \prod_{i=1}^n P(t_i | t_{i-1})$
- The result of these assumptions: $\hat{t} \approx \underset{t_1^n}{\operatorname{argmax}} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$
- Note: B & E represent tag before/after sentence
- HMM taggers are fast and achieve accuracy scores of about 93-95%



Estimating Probability of

- We assume that: $\hat{t} \approx \underset{t_1^n}{\operatorname{argmax}} \prod_{i=1}^n P(w_i|t_i)P(t_i|t_{i-1})$
- Acquire frequencies from a training corpus:
 - Word Frequency with given POS
 - suppose **book** occurs 14 times in a corpus: 10 times (.001) as **NN** (there are 10000 instances of **NN** in the corpus); 3 times (.003) as **VBP** (the corpus has 1000 **VBPs**), and 1 instance of book (.005) as **VB** (the corpus has 500 **VBs**).
 - Given the previous tag, how often does each tag occur
 - suppose **DT is** followed by **NN** 80,000 times (.53), **JJ** 30,000 times (.2), **NNS** 20,000 times (.13), **VCN** 3,000 (.02) times, ... out of a total of 150,000 occurrences of **DT**
- All possible tags for sequence:
 - *The/DT book/{NN|VB|VBP} is/VBZ on/IN the/DT table/{NN|VB|VBP}*
- Hypothetical probabilities for highest scoring tag sequence:
 - *The/DT book/NN is/VBZ on/IN the/DT table/NN*
 - *The/DT=.4, book/NN=.001, is/VBZ=.02, on/IN=.1, the/DT=.4, table/NN=.0005,*
 - **B DT = .61, DT NN = .53, NN VBZ = .44, VBZ IN = .12, IN DT = .05, DT NN = .53 NN E .31**
 - $\prod_{i=1}^n P(w_i|t_i)P(t_i|t_{i-1}) = (.4 \times .61)(.001 \times .53)(.02 \times .44)(.1 \times .12)(.4 \times .05)(.0005 \times .53)(1 \times .31) \approx 2.4 \times 10^{-13}$

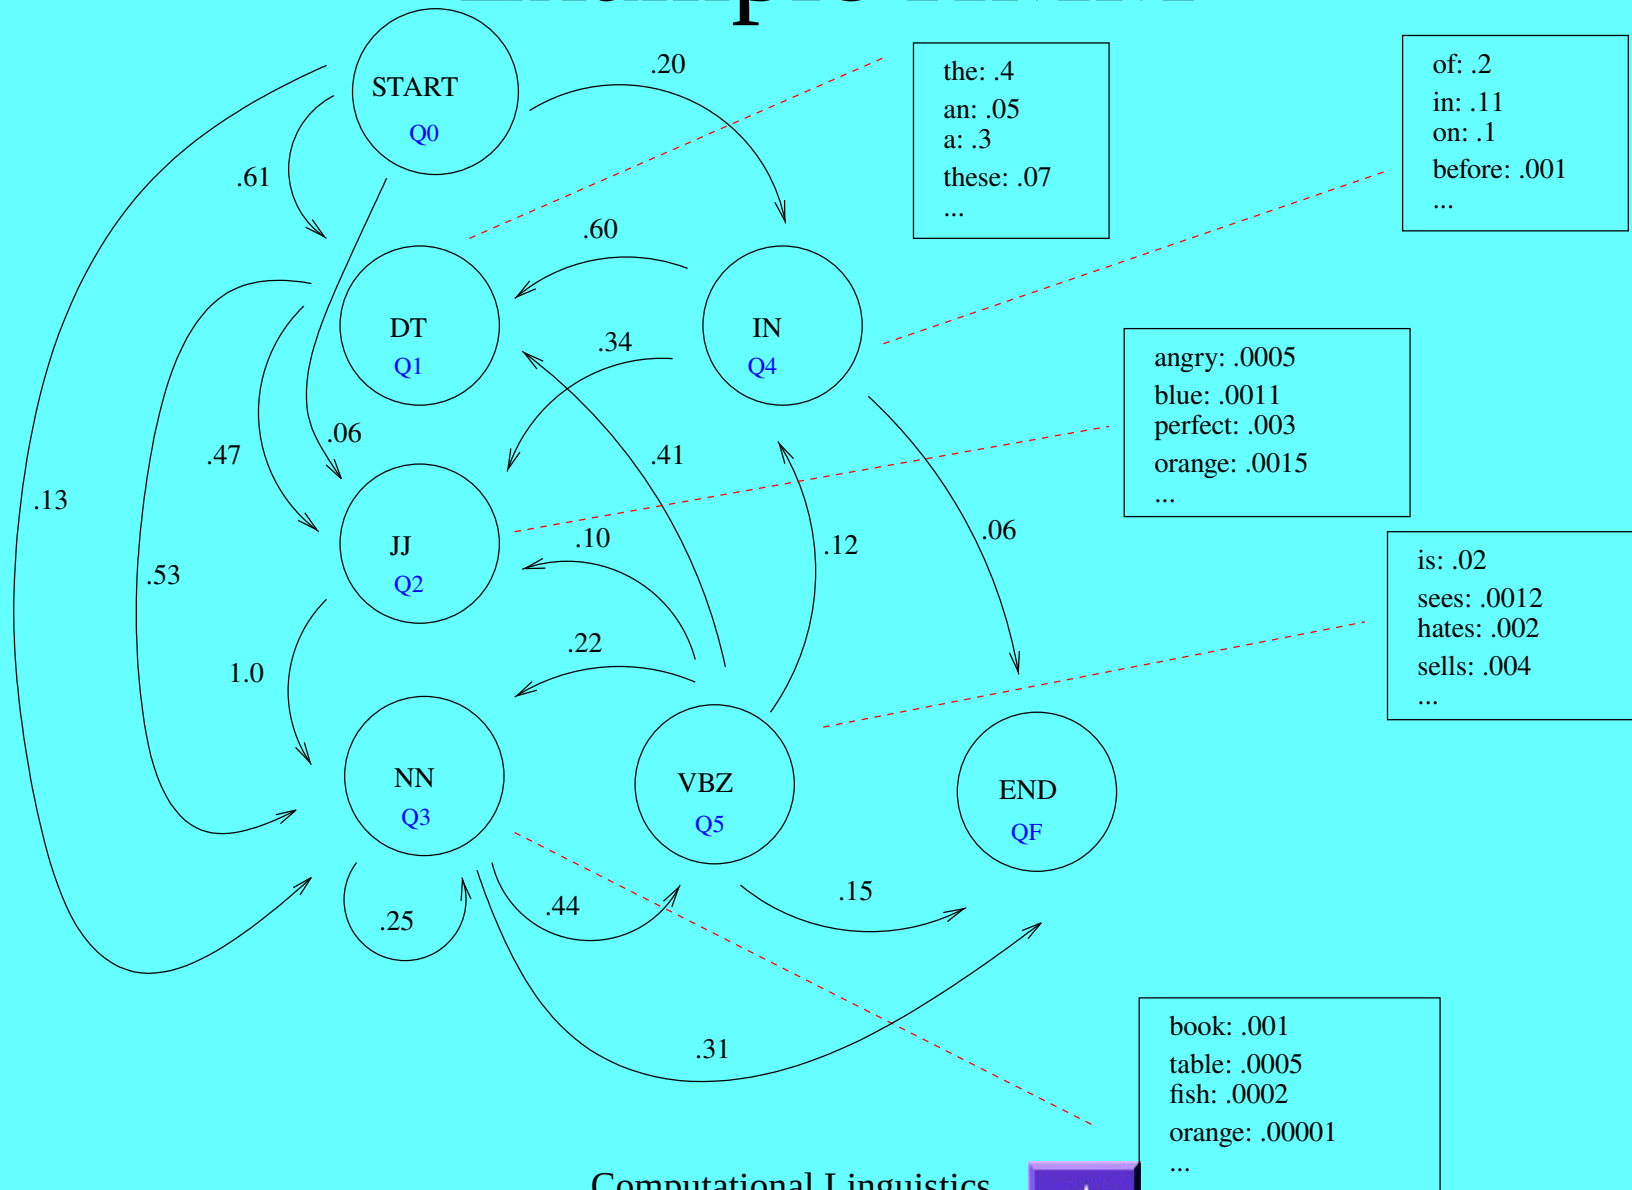


Defining an HMM

- A Weighted Finite-state Automaton (WFSA)
 - Each transition arc is associated with a probability
 - The sum of all arcs outgoing from a single node is 1
- Markov chain is a WFSA in which an input string uniquely determine path through the Automaton
- Hidden Markov Model (HMM) is a slightly different case because some information (previous POS tags) is unknown (or hidden)
- HMM consists of the following:
 - \mathbf{Q} = set of states: \mathbf{q}_0 (start state), ..., \mathbf{q}_F (final state)
 - \mathbf{A} = transition probability matrix of $\mathbf{n} \times \mathbf{n}$ probabilities of transitioning between any pair of \mathbf{n} states ($\mathbf{n} = F+1$). Called: **prior probability** or **transition probability** of a tag sequence
 - \mathbf{O} = sequence of \mathbf{T} observations (**words**) from a vocabulary \mathbf{V}
 - \mathbf{B} = sequence of observation likelihoods (probability of observation generated at state) – Called **likelihood** (of word sequence given tag sequence), aka **emission probability**



Example HMM



Go to Ralph's Viterbi Demo
for *Fish Sleep*



Viterbi Algorithm for HMM

Observed_Words = $w_1 \dots w_T$

- States = $q_0, q_1 \dots q_N, q_F$

$\mathbf{A} = \mathbf{N} \times \mathbf{N}$ matrix such that $a_{i,j}$ is the probability of the transition from q_i to q_j

\mathbf{B} = lookup table such that $b_i(w_t)$ is the probability that POS i is realized as word t

viterbi = $(\mathbf{N}+2) \times \mathbf{T}$ matrix # columns are states, rows are words

backpointer = $(\mathbf{N}+2) \times \mathbf{T}$ matrix # highest scoring previous cells for viterbi

for states q from 1 to \mathbf{N} : ## **BEGINNING**

initialize **viterbi**[$q,1$] to $a_{0,q} * b_q(w_1)$ # score transition $0 \rightarrow q$ given w_1

initialize **backpointer**[$q,1$] to 0 (start state)

for word w from 2 to \mathbf{T} : ## **Middle**

for state q from 1 to \mathbf{N} :

for $\mathbf{T}-1 \times \mathbf{N}$ (w,q) pairs

viterbi[q,w] $\leftarrow \max_{q'=1}^N \text{viterbi}[q',t-1] * a_{q',q} * b_q(w_t)$ # score = maximum previous * prior * likelihood

• **backpointer**[q,w] $\leftarrow \max_{q=1}^N \text{viterbi}[q,T] * a_{q,qF}$ # backpointer = maximum previous

viterbi[qF,T] \leftarrow

END score = maximum previous * prior * likelihood

backpointer[qF,T] \leftarrow

backpointer = maximum previous

- return(best_path) # derive by following backpointers from (qF,T) to q_0



Walk Through: The orange is on the table. (ignoring period)

$$1 * .4 * .61 * .00001 * .53 * .02 * .33 * .1 * .12 * .4 * .6 * .54 * .0005 * .33 * 1 = 2.19 * 10^{-15}$$

	0	1: The	2:orange	3: is	4: on	5: the	6: table	7
Start	1							
DT		.4 * .61				.4*.6		
IN					.1*.12			
JJ			.0015 * .47					
VBZ				.02*.44				
NN			.00001 * .53				.53*.0005	
End								.33 * 1



Comments on Viterbi Trace

- The transitions from B to *the* and the transition from period (.) to E are necessary parts of the process.
 - Scores for first column: transitions from 0 to each possible state given: *the*
 - The probability of reaching Q1 matching the first item on the tape (*the*) will be $.4 \times .61 = .244$ (this is also the only possibility)
 - The transition from period (.) to E will have a high probability.
 - Likelihood that period (.) has POS . is very high (possibly 100%).
- The word *orange* is more likely to be JJ than NN
 - But the path in the chart **B DT JJ** leads to a dead end



Unknown (OOV) Words

- Possibility 1
 - Assume all POS tags have the same probability (e.g., 1/1000)
 - In effect, only use transitions to predict the correct tag
- Possibility 2
 - Use morphology (prefixes, suffixes), orthography (uppercase/lowercase), hyphenation
- Possibility 3:
 - Words occurring once in corpus = instances of UNKNOWN_WORD
 - Distribution of UNKNOWN_WORD used for OOV words
- Possibility 4: Some combination
 - Example: divide UNKNOWN_WORD into morphological classes like UNKNOWN_WORD_ENDING_IN_S



Homework

- <http://cs.nyu.edu/courses/spring23/CSCI-UA.0480-057/homework3.html>
- **Guidance on Program – Next few slides**



Implement Simple version of training stage first

- Data 2 fields (separated by tab): word and POS
 - Start of file = begin of sentence
 - Blank line = begin and end of sentence
 - End of file = end of sentence
- Make 2 hash tables of hash tables (e.g., Python dictionaries of dictionaries)
 1. POS → table of frequencies of words that occur with that POS
 - Example: likelihood['DT'] → {'the':1500,'a':200,'an':100, ...}
 - Hash table of POSs with each value a hash table from words to frequencies
 2. STATE → table of frequencies of following states
 - Example: Transition['Begin_Sent'] → {'DT':1000,'NNP':500,'VB':200, ...}
 - Example: Transition['DT'] → {'NN':500,'NNP':200,'VB':30,...}
 - Hash table of states with each a value a hash table from states to frequencies
 - States = Begin_Sent, End_Sent and all POSs
 - Alternative to dictionary of dictionaries: 2 dimensional array:
 - Each dimension should be 2 + number of different POS tags
- Go through the data one line at a time
 - Record frequencies for both 1 and 2
 - Loop thru hash table and convert frequencies into probabilities
 - $\text{freq}/\text{total} = \text{probability}$



Simple Version of Transducer

- Make a 2 dimensional array (or equivalent)
 - columns represent tokens at positions in the text
 - 0 = start of sentence
 - N = N th token (word punctuation at position N)
 - $\text{Length}+1$ = end of sentence
 - rows represent S states: the start symbol, the end symbol and all possible POS (NN, JJ, ...)
 - cells represent the likelihood that a particular word is at a particular state
- Traverse the chart as per the algorithm (fish sleep slides, etc.)
 - For all states at position 1, multiply transition probability from Start (position 0) by likelihood that word at position 1 occurs in that state. Choose highest score for each cell.
 - For n from 2 to N (columns)
 - for each cell $[n,s]$ in column n and each state $[n-1,s']$ in column $n-1$:
 - get the product of:
 - likelihood that token n occurs in state s
 - the transition probability from s' to s
 - the score stored in $[n-1,s']$
 - At each position $[n,s]$, record the max of the s scores calculated



Calculating Probabilities

- The probability of each transition to state N for token T is assumed to be the product of 3 factors
 - Probability that state N occurs with token T
 - There is 100% chance that the start state will be at the beginning of the sentence
 - There is 100% chance that the end state will be at the end of the sentence
 - If a token was observed in the training corpus, look up probability from table
 - For Out of Vocabulary words, there are several strategies
 - Simple strategy (for first implementation): 1/1000 or 100% divided by number of states or any fraction that is the same for all POS
 - Other strategies are a separate discussion
 - Probability that state N occurs given previous state
 - Look up in table, calculate for every possible previous state
 - Highest Probability of previous state (calculate for each previous state)
- For each new state, choose the highest score (this is the bigram model)
- Choose the POS tag sequence resulting in the highest score in the end state



OOV Strategies from slide 21

- Default (use until other parts of program are debugged)
 - Assume all POS tags have the same probability (e.g., 1/1000)
 - In effect, only use transitions to predict the correct tag
- Morphology
 - Use prefixes, suffixes, uppercase/lowercase, hyphenation, to predict POS classes of OOV words
 - Assign “made up” values based on these features?
 - Perhaps hard-code unusual punctuation based on Penn Treebank specs
- Compute probability of UNKNOWN_WORD
 - Treat words occurring once in training collectively as UNKNOWN_WORD
 - don't record them separately (recalculate likelihood table)
 - UNKNOWN_WORD probability used for OOV words by transducer
- Combination:
 - UNKNOWN_ending_in_s, UNKNOWN_ending_in_ed, UNKNOWN_with_capital_letter, ...



How you Might Improve your Score

- Do error analysis on development corpus – base changes on your findings
 - Example: Are there errors for punctuation (which should be almost unambiguous)?
- Implement a trigram algorithm
 - See Jurafsky and Martin (p. 149)
 - 4-gram is a waste of time for this size corpus
- Note: A clever OOV system contributes more to score than trigram
- Manual rule system using constraints, e.g., slide 7.
 - For words with frequency > 1, assume the disjunction of observed labels is possible
 - Rule out possibilities according to constraints
 - Run this and compare results with HMM system
 - Figure out way of combining results with HMM based on error analysis
 - Voting, weighted combinations, etc.



Grading

- Your grade 1–11 is based on:
 - Meeting the format constraints (1 point)
 - Your accuracy score
 - 95% → 10 points, 93% → 9 points
 - 92% → 8 points, 91% → 7 points
 - 90% → 6 points, 85% → 5 points
- You should include a short write-up of what you did, so it is easier to evaluate.

